

Artificial Intelligence



Lesson 1

About

- Lecturer: Prof. Sarit Kraus
- TA: Galit Haim: haimga@cs.biu.ac.il
- (almost) All you need can be found on the course website:
 - <http://u.cs.biu.ac.il/~haimga/Teaching/AI/>

Course Requirements 1

- The grade is comprised of 70% exam and 30% exercises.
- 3 programming exercises will be given. Work individually.
- All the exercises are counted for the final grade.
- 10% each.
- Exercises will be written in C++ or JAVA only. It should be compiled and run on planet machine, and will be submitted via “submit”. Be precise!

Course Requirements 2

- Exercises are not hard, but work is required. Plan your time ahead!
- When sending me mail **please include the course number (89-570) in the header**, to pass the automatic spam filter.
- You (probably) will be required to participate in AI experiments.
- See other general rules in:
<http://u.cs.biu.ac.il/~haimga/Teaching/AI/assignments/general-rules.pdf>

Course Schedule

- Lesson 1:
 - Introduction
 - Transferring a general problem to a graph search problem.
- Lesson 2
 - Uninformed Search (BFS, DFS etc.).
- Lesson 3
 - Informed Search (A*, Best-First-Search etc.).

Course Schedule – Cont.

- Lesson 4
 - Local Search (Hill Climbing, Genetic algorithms etc.).
- Lesson 5
 - “Search algorithms” chapter summary.
- Lesson 6-7
 - Game-Trees: Min-Max & Alpha-Beta algorithms.

Course Schedule – Cont.

- Lesson 8-9
 - Planning: STRIPS algorithm
- Lesson 10-11-12
 - Learning: Decision-Trees, Neural Network, Naïve Bayes, Bayesian Networks and more.
- Lesson 13: Robotics
- Lesson 14
 - Questions and exercise.

AI – Alternative Definitions

- **Elaine Rich and Kevin Knight:** AI is the study of how to make computers do things at which, at the moment, people are better.
- **Stuart Russell and Peter Norvig:** [AI] has to do with smart programs, so let's get on and write some.
- **Claudson Bornstein:** AI is the science of common sense.
- **Douglas Baker:** AI is the attempt to make computers do what people think computers cannot do.
- **Astro Teller:** AI is the attempt to make computers do what they do in the movies.

AI Domains

- Games – chess, checkers, tile puzzle.
- Expert systems
- Speech recognition and Natural language processing, Computer vision, Robotics.

AI & Search

- "The two most fundamental concerns of AI researchers are *knowledge representation* and *search*"
- "*knowledge representation* ... addresses the problem of capturing in a language...suitable for computer manipulation"
- "*Search* is a problem-solving technique that systematically explores a space of problem states".Luger, G.F. Artificial Intelligence: Structures and Strategies for Complex Problem Solving

Solving Problems with Search Algorithms

- Input: a problem P .
- Preprocessing:
 - Define *states* and a *state space*
 - Define *Operators*
 - Define a *start* state and *goal* set of states.
- Processing:
 - Activate a Search algorithm to find a *path* from start to one of the goal states.

Example - Missionaries & Cannibals

- State space – $[M,C,B]$
- Initial State – $[3,3,1]$
- Goal State – $[0,0,0]$
- Operators – adding or subtracting the vectors $[1,0,1]$, $[2,0,1]$, $[0,1,1]$, $[0,2,1]$ or $[1,1,1]$
- Path – moves from $[3,3,1]$ to $[0,0,0]$
- Path Cost – river trips
- <http://www.plastelina.net/game2.html>
- <http://www.youtube.com/watch?v=W9NEWxabGmg>

Breadth-First-Search Pseudo code

- Intuition: Treating the graph as a tree and scanning top-down.
- Algorithm:

BFS(Graph graph, Node start, Vector Goals)

1. $L \rightarrow \text{make_queue}(\text{start})$
2. While L not empty loop
 1. $n \leftarrow L.\text{remove_front}()$
 2. If goal (n) return true
 3. $S \leftarrow \text{successors}(n)$
 4. $L.\text{insert}(S)$
3. Return false

Breadth-First-Search Attributes

- Completeness – yes ($b < \infty, d < \infty$)
- Optimality – yes, if graph is unweighted.
- Time Complexity: $O(b^{d+1})$
- Memory Complexity: $O(b^{d+1})$
 - Where b is branching factor and d is the solution depth
- See water tanks example:
- http://u.cs.biu.ac.il/~haimga/Teaching/AI/lessons/lesson1&2b-%20WaterTank-BFS_DFS.pdf

Artificial Intelligence



Lesson 2

Uninformed Search

- Uninformed search methods use only information available in the problem definition.
 - Breadth First Search (BFS)
 - Depth First Search (DFS)
 - Iterative DFS (IDA)
 - Bi-directional search
 - Uniform Cost Search (a.k.a. Dijkstra alg.)


Depth-First-Search Pseudo code

DFS(Graph graph, Node start, Vector Goals)

1. $L \leftarrow \text{make_stack}(\text{start})$
2. While L not empty loop
 - 2.1 $n \leftarrow L.\text{remove_front}()$
 - 2.2 If goal (n) return true
 - 2.3 $S \leftarrow \text{successors}(n)$
 - 2.4 $L.\text{insert}(S)$
3. Return false

Depth-First-Search Attributes

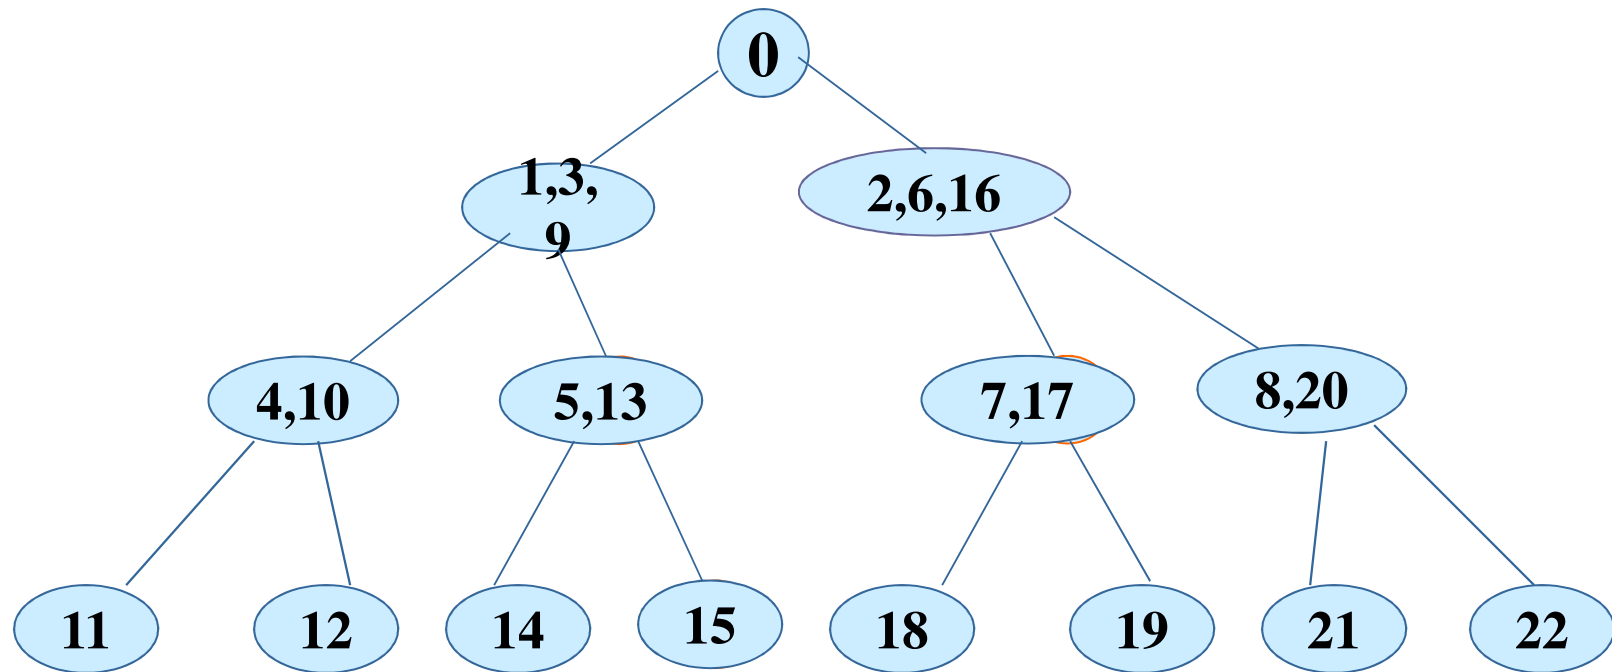
- Completeness – No. Infinite loops or Infinite depth can occur.
- Optimality – No.
- Time Complexity: $O(b^m)$
- Memory Complexity: $O(bm)$
 - Where b is branching factor and m is the maximum depth of search tree
- See water tanks example



Limited DFS Attributes

- Completeness – Yes, if $d \leq l$
- Optimality – No.
- Time Complexity: $O(b^l)$
 - If $d < l$, it is larger than in BFS
- Memory Complexity: $O(bl)$
 - Where b is branching factor and l is the depth limit.

Depth-First Iterative-Deepening



The numbers represent the order generated by DFID

Iterative-Deepening Attributes


- Completeness – Yes
- Optimality – yes, if graph is un-weighted.
- Time Complexity:

$$O((d)b + (d-1)b^2 + \dots + (1)b^d) = O(b^d)$$

- Memory Complexity: $O(db)$
 - Where b is branching factor and d is the maximum depth of search tree


State Redundancies

- [illegible]



Bi-directional Search

- Search both from initial state to goal state.
- Operators must be symmetric.



Bi-directional Search Attributes

- Completeness – Yes, if both directions use BFS
- Optimality – yes, if graph is un-weighted and both directions use BFS.
- Time and memory Complexity: $O(b^{d/2})$
- Pros.
 - Cuts the search tree by half (at least theoretically).
- Cons.
 - Frontiers must be constantly compared.

Minimum cost path


- General minimum cost path-search problem:
 - Find shortest path from start state to one of the goal states in a weighted graph.
 - Path cost function is $g(n)$: sum of weights from start state to goal.

Uniform Cost Search

- Also known as Dijkstra's algorithm.
- Expand the node with the minimum path cost first.
- Implementation: priority queue.

Example of Uniform Cost Search


- Assume an example tree with different edge costs, represented by numbers next to the edges.




Notations for this example:

- generated node
- expanded node

Example of Uniform Cost Search





Closed list:




Open list:

28

Example of Uniform Cost Search





Closed list:




Open list:

Example of Uniform Cost Search




Closed list:




Open list:


Example of Uniform Cost Search




Closed list:




Open list:




Example of Uniform Cost Search




Closed list:




Open list:




Example of Uniform Cost Search





Closed list:



Open list:



Example of Uniform Cost Search





Closed list:



Open list:

Example of Uniform Cost Search



Closed list:

Open list:

Uniform Cost Search Attributes

- Completeness: yes, for positive weights
- Optimality: yes
- Time & Memory complexity: $O(b^{\lfloor c/e \rfloor})$
 - Where b is branching factor, c is the optimal solution cost and e is the minimum edge cost (each operator cost at least e).
 - UCS uses cost path and not length path, therefore, the ‘ d ’ term is not used here.

Informed Search

- Incorporate additional measure of a potential of a specific state to reach the goal.
- A potential of a state to reach a goal is measured through a heuristic function $h(n)$.
- An evaluation function is denoted $f(n)$.

Best First Search Algorithms

- Principle: Expand node n with the best evaluation function value $f(n)$.
- Implement via a priority queue
- Algorithms differ with definition of f :
 - Greedy Search: $f(n) = h(n)$
 - A^* : $f(n) = g(n) + h(n)$
 - IDA*: iterative deepening version of A^*
 - Etc'

Exercise

- Q: Does a Uniform-Cost search be considered as a Best-First algorithm?
- A: Yes. It can be considered as a Best-First algorithm with evaluation function $f(n)=g(n)$.
- Q: In what scenarios IDS outperforms DFS?, BFS?
- A:
 - IDS outperforms DFS when the search tree is a lot deeper than the solution depth.
 - IDS outperforms BFS when BFS run out of memory.

Exercise – Cont.

- Q: Why do we need a closed list?
- A: Generally a closed list has two main functionalities:
 - Prevent re-exploring of nodes.
 - Hold solution path from start to goal (DFS based algorithms have it anyway).
- Q: Does Breadth-FS find optimal path length in general?
- A: No, unless the search graph is un-weighted.
- Q: Will IDS always find the same solution as BFS given that the nodes expansion order is deterministic?
- A: Yes. Each iteration of IDS explores new nodes the same order a BFS does.